

nimbleHMC: An R package for Hamiltonian Monte Carlo sampling in nimble

13 August 2022

Summary

Markov chain Monte Carlo (MCMC) algorithms are widely used for fitting hierarchical (graphical) models to observed data, and more generally, for simulating from high-dimensional probability distributions. MCMC is the predominant tool used in Bayesian analyses, where the distribution of interest (the “target distribution”) is defined as the posterior distribution of the unknown model parameters conditional on the data. MCMC does not specify a single algorithm, but rather a family of algorithms admitting any assignment of valid sampling techniques (“samplers”) to the unobserved model dimensions. There exist a vast and diverse landscape of valid samplers to draw upon, which differ significantly in their underlying approaches to the sampling problem, complexity, autocorrelation of the samples produced, and applicability.

Hamiltonian Monte Carlo (HMC; Brooks et al. 2011) is one such sampling technique which can be applied to any subset of continuous-valued model dimensions. HMC uses the gradient of the target distribution to generate large transitions (in parameter space) in the output sequence of samples. This results in low autocorrelation, and therefore high information content. That is, the samples generated using HMC are more likely to be highly informative about the target distribution of interest, relative for example to an equal-length sequence of highly autocorrelated samples. This rich information content does not come freely, however, as calculating gradients of the target distribution is computationally expensive.

There exist numerous software packages which provide implementations of MCMC for mainstream use such as **nimble** (Valpine et al. 2017), **jags** (Plummer and others 2003), **pyMC** (Fonnesbeck et al. 2015), and **Stan** (Carpenter et al. 2017). Each such package provides a language for specifying general hierarchical model structures, and supplying data. Following specification of the problem, each package generates an MCMC algorithm which specifically samples from the target posterior distribution of the specified model, and executes this algorithm to generate a sequence of samples from this distribution. These packages differ, however, in their approaches to sampler assignment for each unobserved model dimension. As sampling techniques vary in terms of computational demands and the quality of the samples produced, the effectiveness of the MCMC algorithms may vary depending on the software used, and the particular model at hand. Each software package provides a valid, but distinct approach for assigning samplers to define the MCMC algorithm.

Among general-purpose MCMC software packages, **nimble** uniquely provides the ability to specify which samplers are applied to each model dimension. Prior to generating an executable MCMC algorithm, **nimble** has the intermediate stage of MCMC configuration. At configuration time, users may select any valid assignment of samplers to each unobserved model dimension, mixing and matching between those samplers provided with **nimble**. The base **nimble** package provides a variety of non-derivative-based samplers, including random walk Metropolis-Hastings (Robert and Casella 1999), slice sampling (Neal 2003), conjugate samplers (George, Makov, and Smith 1993), and many others. After configuration is finished, an MCMC algorithm is generated according to the sampler assignments therein, and executed to generate a sequence of samples.

The **nimbleHMC** package provides an implementation of HMC sampling which is compatible for use within **nimble**. Specifically, **nimbleHMC** implements the No-U-Turn variety of HMC (HMC-NUTS; Hoffman, Gelman, and others 2014), which removes the necessity of hand-specifying tuning parameters of the HMC sampler. Using **nimbleHMC**, HMC samplers can be assigned to any subset of continuous-valued model dimensions at the time of **nimble**’s MCMC configuration, which may be used in combination with any other samplers provided with the base **nimble** package.

Statement of need

HMC is recognized as a state-of-the-art MCMC sampling strategy, capable of efficiently generating samples with strong inferential power. A testimony to this, packages such as **Stan** have built software exclusively around the use of HMC sampling. As a result, however, such software is unable to operate on models with discrete (non-continuous) valued dimensions, a result of the non-applicability of HMC. Models with discrete-valued dimensions arise in a range of common statistical motifs including hidden Markov models, finite mixture models, and generally in the presence of unobserved categorical data, among others (Bartolucci, Pandolfi, and Pennoni 2022). In contrast, other mainstream MCMC packages such as **WinBUGS**, **OpenBUGS** and **jags** have the ability to sample discrete model dimensions, but do not implement HMC. This leaves a gap, as there is no support for applying HMC sampling to continuous-valued dimensions of hierarchical models which also contain discrete dimensions.

It is an open question what MCMC algorithm, or which combination of samplers, will optimize the fitting of any particular hierarchical model and dataset. The metric of interest is the effective sample size of the sequence of samples (which accounts for autocorrelation) generated per unit runtime of the MCMC. That is, how quickly an MCMC algorithm generates meaningful information to characterize the target distribution. This metric is called MCMC efficiency (Turek et al. 2017), but what assignment of samplers maximizes this metric is a difficult and open question (Ponisio et al. 2020). For that reason, the ability to mix-and-match samplers from among as large a pool of candidates as possible is important from both practical and theoretical standpoints. Indeed, there even exist packages such as **compareMCMCs** (Valpine, Paganin, and Turek 2022), the purpose of which is to compare the relative performance of distinct MCMC algorithms.

The **nimble** package uniquely provides the ability to custom-specify the assignment of MCMC sampling algorithms, which allows the exploration and the study of efficient approaches to MCMC. No existing software to date can operate on discrete model dimensions *and* offers the option of HMC sampling. Here, the **nimbleHMC** package augments the **nimble** package by providing an HMC sampler suitable for use within **nimble**'s MCMC. This fills the gap, allowing HMC samplers to operate alongside the existing continuous and discrete sampling algorithms available in **nimble**.

References

- Bartolucci, Francesco, Silvia Pandolfi, and Fulvia Pennoni. 2022. "Discrete Latent Variable Models." *Annual Review of Statistics and Its Application* 9: 425–52.
- Brooks, Steve, Andrew Gelman, Galin Jones, and Xiao-Li Meng. 2011. *Handbook of Markov Chain Monte Carlo*. CRC press.
- Carpenter, Bob, Andrew Gelman, Matthew D Hoffman, Daniel Lee, Ben Goodrich, Michael Betancourt, Marcus Brubaker, Jiqiang Guo, Peter Li, and Allen Riddell. 2017. "Stan: A Probabilistic Programming Language." *Journal of Statistical Software* 76 (1).
- Fonnesbeck, Chris, Anand Patil, David Huard, and John Salvatier. 2015. "PyMC: Bayesian Stochastic Modelling in Python." *Astrophysics Source Code Library*, ascl-1506.
- George, Edward I, UE Makov, and AFM Smith. 1993. "Conjugate Likelihood Distributions." *Scandinavian Journal of Statistics*, 147–56.
- Hoffman, Matthew D, Andrew Gelman, and others. 2014. "The No-U-Turn Sampler: Adaptively Setting Path Lengths in Hamiltonian Monte Carlo." *J. Mach. Learn. Res.* 15 (1): 1593–1623.
- Neal, Radford M. 2003. "Slice Sampling." *The Annals of Statistics* 31 (3): 705–67.
- Plummer, Martyn, and others. 2003. "JAGS: A Program for Analysis of Bayesian Graphical Models Using Gibbs Sampling." In *Proceedings of the 3rd International Workshop on Distributed Statistical Computing*, 124:1–10. 125.10. Vienna, Austria.
- Ponisio, Lauren C, Perry de Valpine, Nicholas Michaud, and Daniel Turek. 2020. "One Size Does Not Fit

All: Customizing Mcmc Methods for Hierarchical Models Using Nimble.” *Ecology and Evolution* 10 (5): 2385–2416.

Robert, Christian P, and George Casella. 1999. “The Metropolis—Hastings Algorithm.” In *Monte Carlo Statistical Methods*, 231–83. Springer.

Turek, Daniel, Perry de Valpine, Christopher J Paciorek, and Clifford Anderson-Bergman. 2017. “Automated Parameter Blocking for Efficient Markov Chain Monte Carlo Sampling.” *Bayesian Analysis* 12 (2): 465–90.

Valpine, Perry de, Sally Paganin, and Daniel Turek. 2022. “CompareMCMCs: An R Package for Studying Mcmc Efficiency.” *Journal of Open Source Software* 7 (69): 3844.

Valpine, Perry de, Daniel Turek, Christopher J Paciorek, Clifford Anderson-Bergman, Duncan Temple Lang, and Rastislav Bodik. 2017. “Programming with Models: Writing Statistical Algorithms for General Model Structures with Nimble.” *Journal of Computational and Graphical Statistics* 26 (2): 403–13.